RHOS 4 + RHEL-HA Integration

Fabio M. Di Nitto

01 March 2014

FAQ

- Why not RabbitMQ?
- Why not Percona or some other replicated DB?

*** we do NOT ship them, we CANNOT support them ***

- Why ha-proxy and not some other random LB?

* we can only test against what we support and have in the house *

The process basics

- Service should be all managed/monitored by pacemaker based cluster, with access via load-balancer for endpoints (no direct connections)
- Service isolation: every service should be capable of running in its own isolated cluster
  - Better identify service inter-dependency (start-up order)
  - Lay down the basic for “extreme” scale solutions
  - Easily identifies requirements per service
  - Spot weakness in RCP/network protocols and documented configuration steps
- Service can collapse from standalone to one cluster
- Can be any random number of nodes between 2 and 16 per given cluster (tested on 2 nodes)

The road so far

- Creating bricks with pacemaker:
  - LB - Mixed
  - Mysql - A/P
  - Qpid - hot stand-by
  - Keystone
  - memcached – hot stand-by
  - Glance
  - Cinder
  - Neutron – A/P
  - Swift
  - Nova
  - Horizon
  - Heat
  - Ceilometer

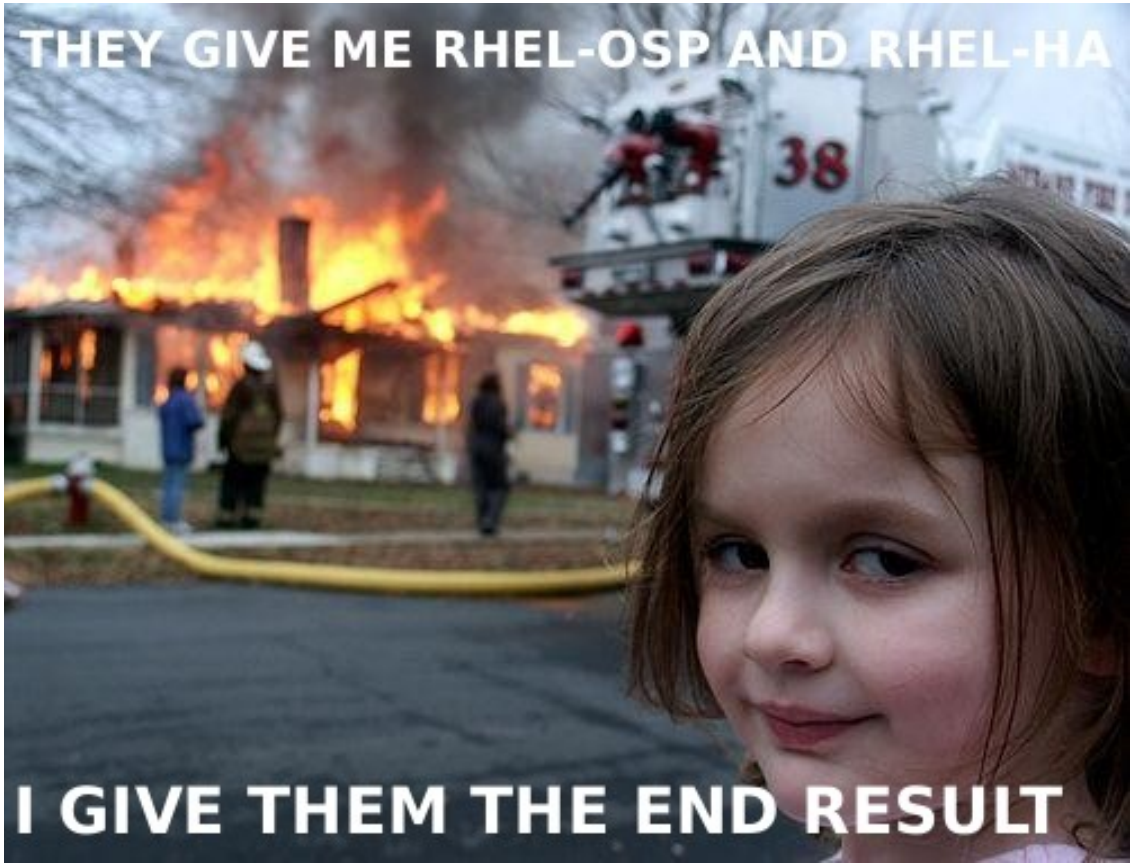

Current deployment overview: the art of too many clusters

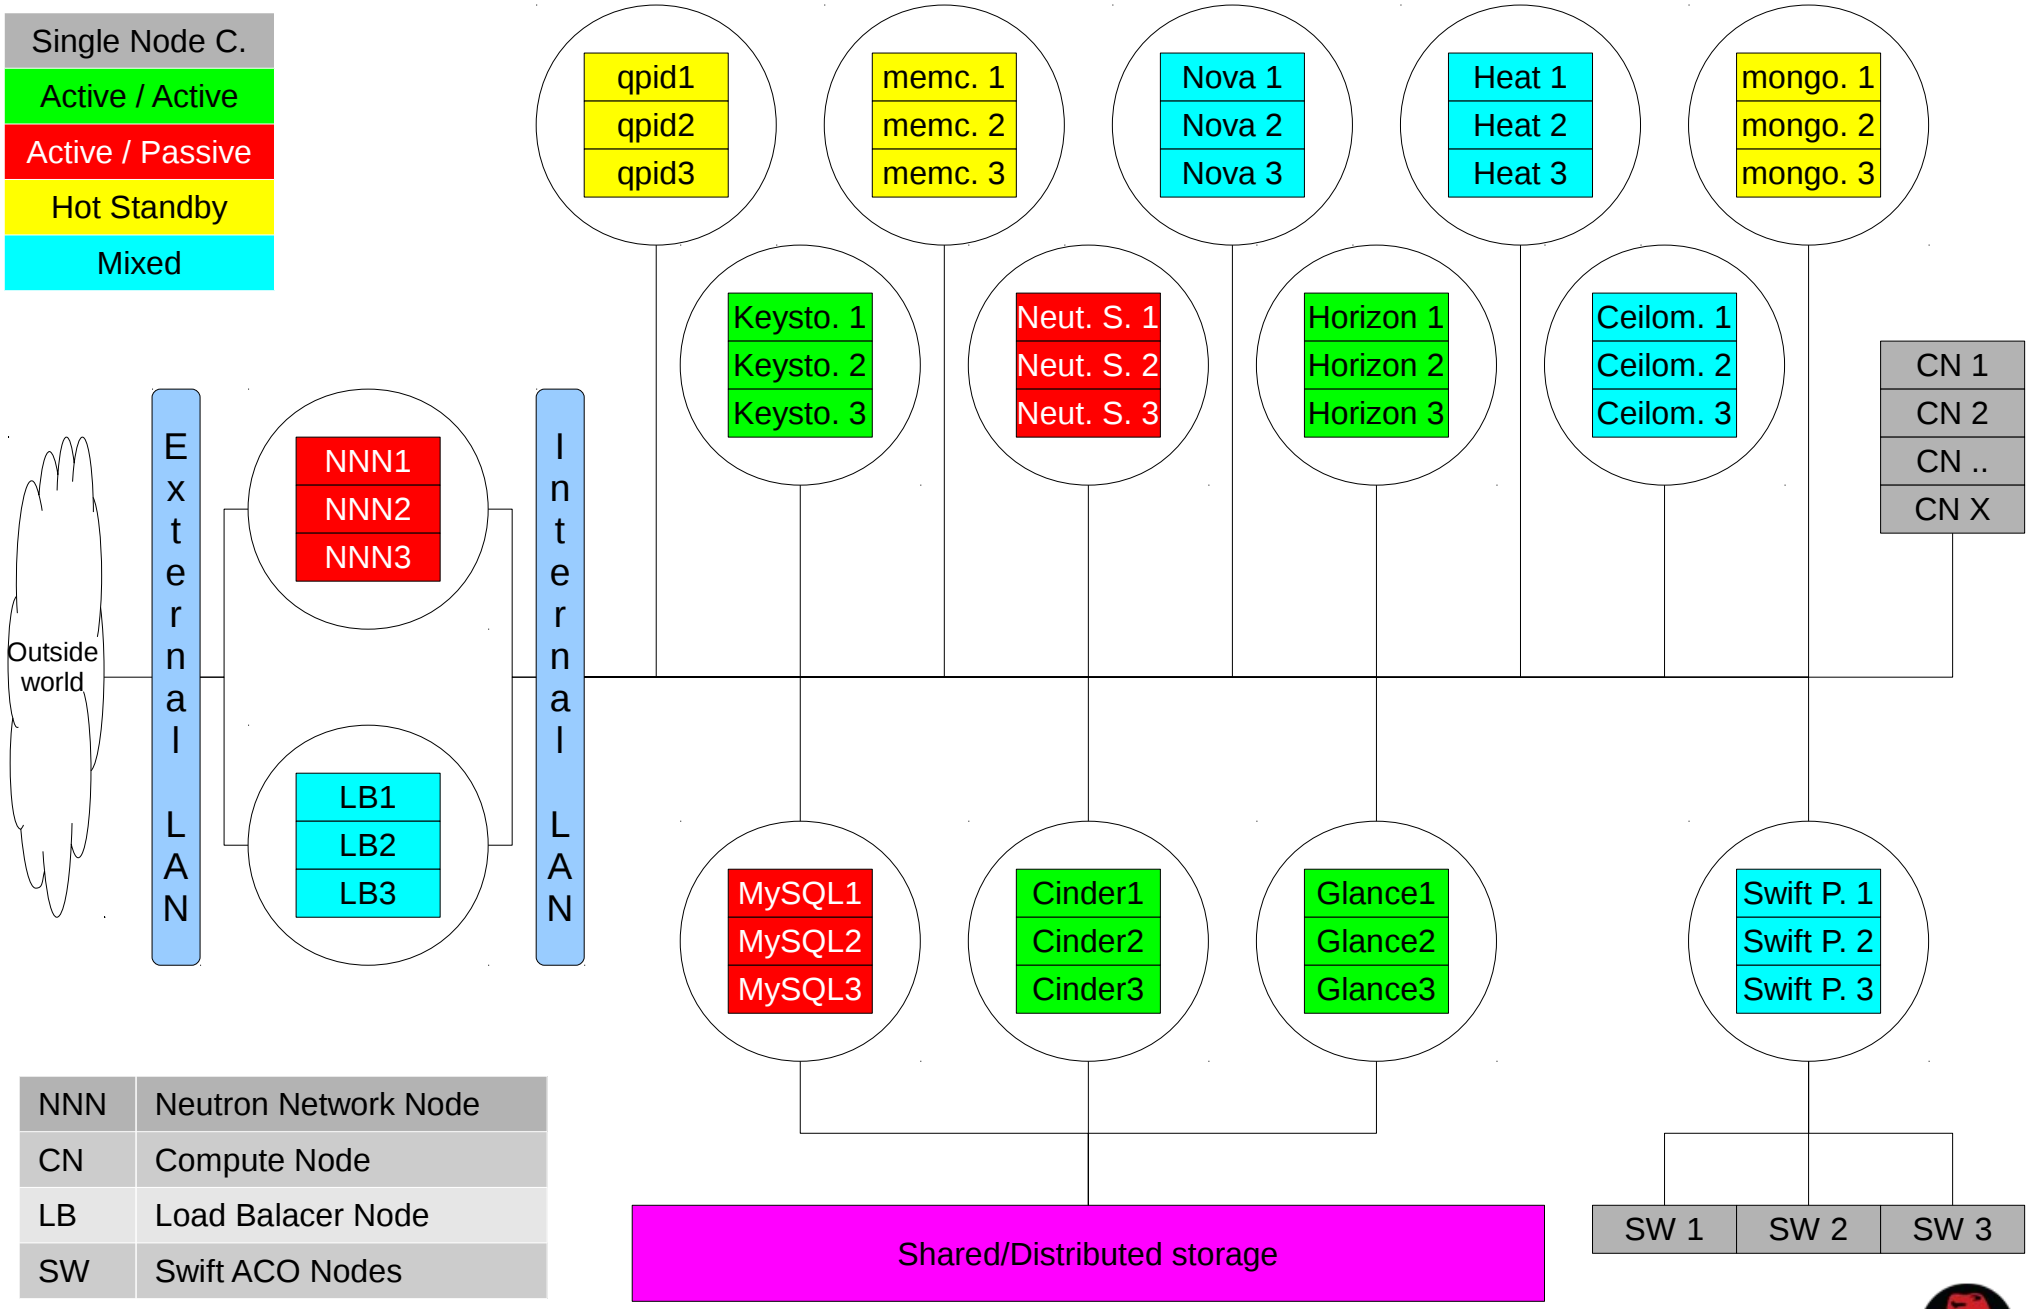

NNN	Neutron Network Node
CN	Compute Node
LB	Load Balacer Node
SW	Swift ACO Nodes

Inter-dependencies and start-up order

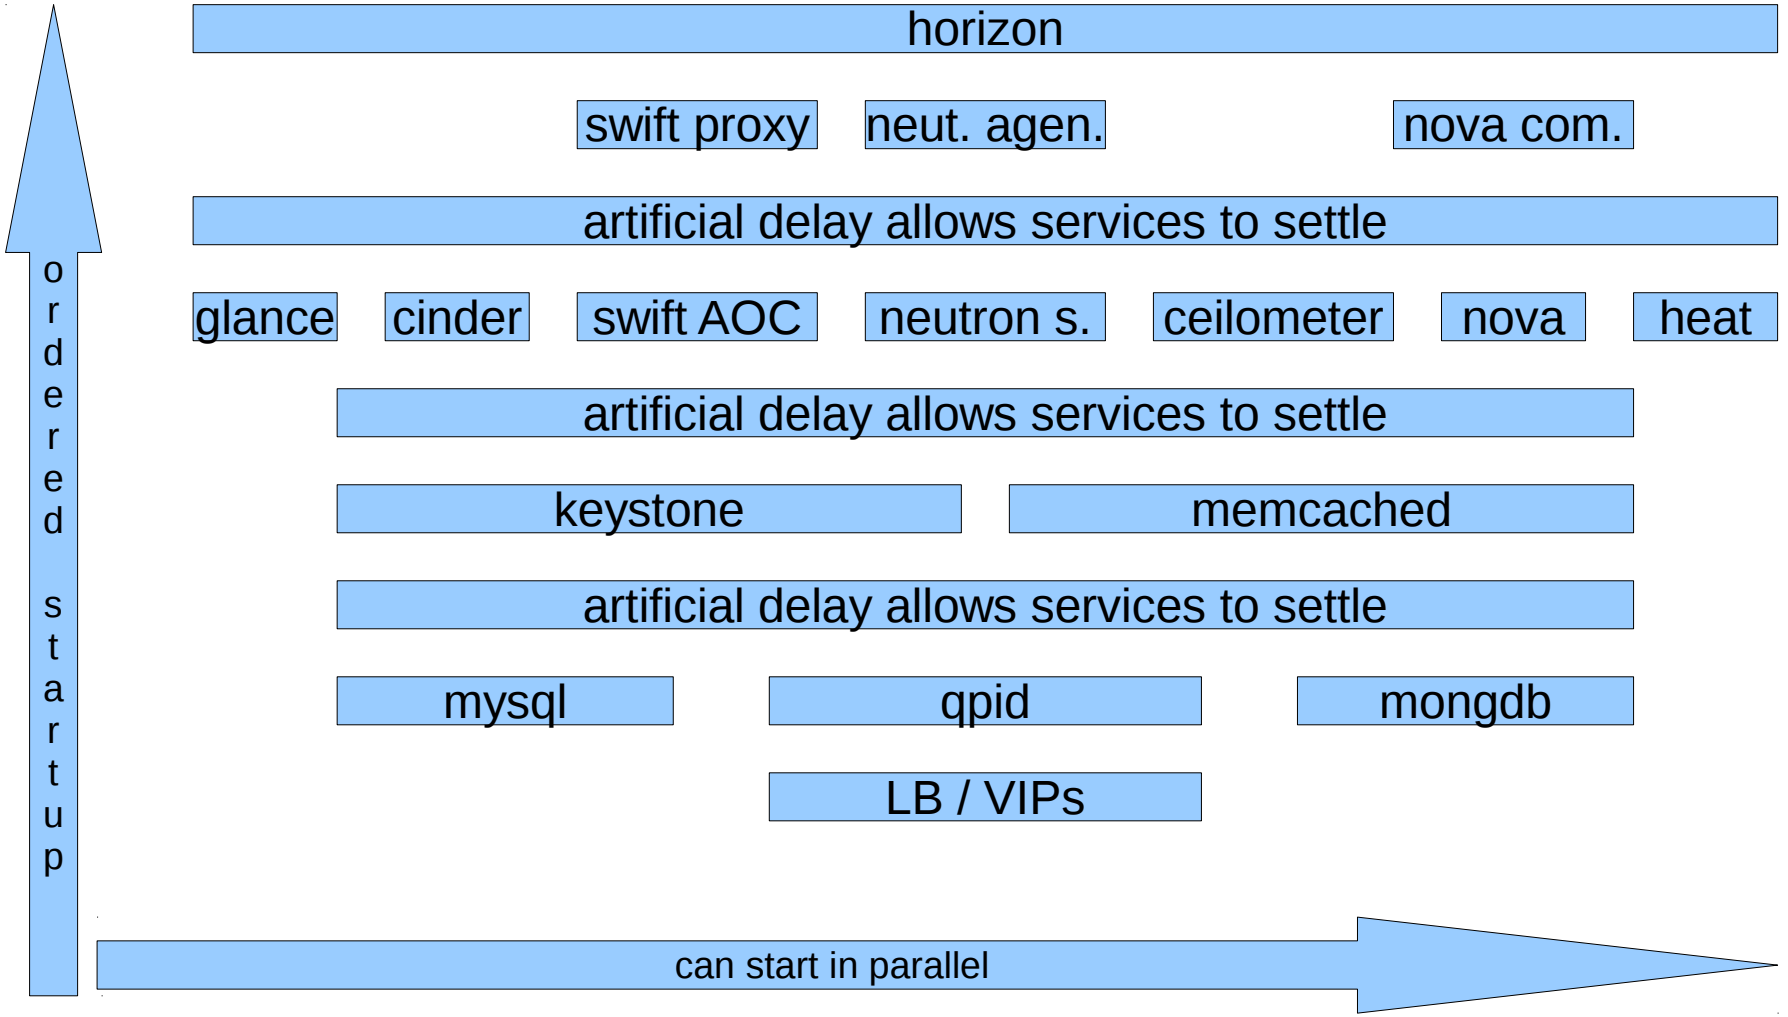

The importance of the

Load Balancer